

Neural Network

J. Hartsfield

March 17, 2017

Neural Network examples using neuralnet package

Neural network with no hidden layers is just regression


```
library(neuralnet)
```

```
## Warning: package 'neuralnet' was built under R version 3.3.3
```

```
load("Auto.rda")
n = length(Auto$mpg)
set.seed(1234)
Z = sample(n, 200)
Auto.train = Auto[Z,]
Auto.test = Auto[-Z,]
nn0 = neuralnet(mpg ~ weight + acceleration + horsepower + cylinders, data=Auto.train, hidden=0)
summary(nn0)
```

| ## | Length | Class | Mode |
|------------------------|--------|------------|----------|
| ## call | 4 | -none- | call |
| ## response | 200 | -none- | numeric |
| ## covariate | 800 | -none- | numeric |
| ## model.list | 2 | -none- | list |
| ## err.fct | 1 | -none- | function |
| ## act.fct | 1 | -none- | function |
| ## linear.output | 1 | -none- | logical |
| ## data | 9 | data.frame | list |
| ## net.result | 1 | -none- | list |
| ## weights | 1 | -none- | list |
| ## startweights | 1 | -none- | list |
| ## generalized.weights | 1 | -none- | list |
| ## result.matrix | 8 | -none- | numeric |

```
plot(nn0, rep="best")
```



Error: 1829.965357 Steps: 6697

Compare this to linear regression:

```
linMod = lm(mpg ~ weight + acceleration + horsepower + cylinders, data=Auto.train)
summary(linMod)
```

```
##
## Call:
## lm(formula = mpg ~ weight + acceleration + horsepower + cylinders,
##     data = Auto.train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -10.5315078  -2.5910769  -0.3724999   2.4736324  16.3477169
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  42.487281370  3.555224343  11.95066 < 0.000000000000000222 ***
## weight      -0.005014105  0.001062895  -4.71741  0.000004546 ***
## acceleration  0.148018374  0.181142033   0.81714   0.414846
## horsepower   -0.050314865  0.022368342  -2.24938   0.025606 *
## cylinders    -0.258574594  0.430455304  -0.60070   0.548737
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.332306 on 195 degrees of freedom
## Multiple R-squared:  0.697537, Adjusted R-squared:  0.6913326
## F-statistic: 112.4267 on 4 and 195 DF, p-value: < 0.00000000000000022204
```

```
# Predicted data from lm
pr.lm <- predict(linMod,Auto.test)
```

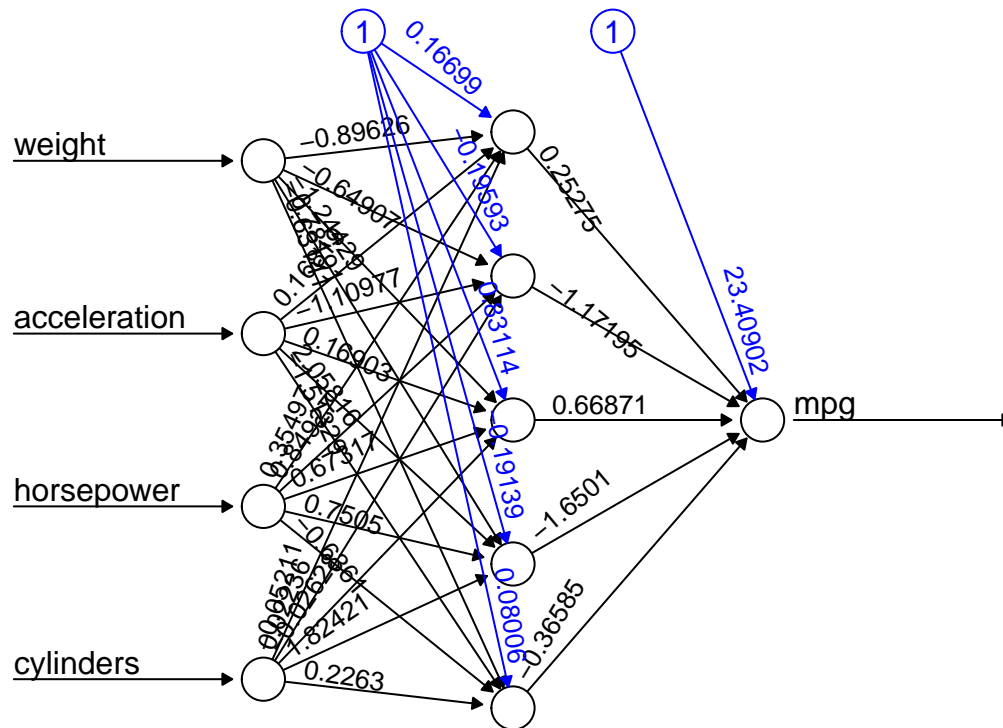
```
# Test MSE
MSE.lm <- sum((pr.lm - Auto.test$mpg)^2)/nrow(Auto.test)
```

Now we introduce a hidden layer that has 5 nodes.

```
nn5 =neuralnet(mpg ~ weight + acceleration + horsepower + cylinders, data=Auto.train, hidden=5)
summary(nn5)
```

| ## | Length | Class | Mode |
|------------------------|--------|------------|----------|
| ## call | 4 | -none- | call |
| ## response | 200 | -none- | numeric |
| ## covariate | 800 | -none- | numeric |
| ## model.list | 2 | -none- | list |
| ## err.fct | 1 | -none- | function |
| ## act.fct | 1 | -none- | function |
| ## linear.output | 1 | -none- | logical |
| ## data | 9 | data.frame | list |
| ## net.result | 1 | -none- | list |
| ## weights | 1 | -none- | list |
| ## startweights | 1 | -none- | list |
| ## generalized.weights | 1 | -none- | list |
| ## result.matrix | 34 | -none- | numeric |

```
#nn5
plot(nn5, rep="best")
```



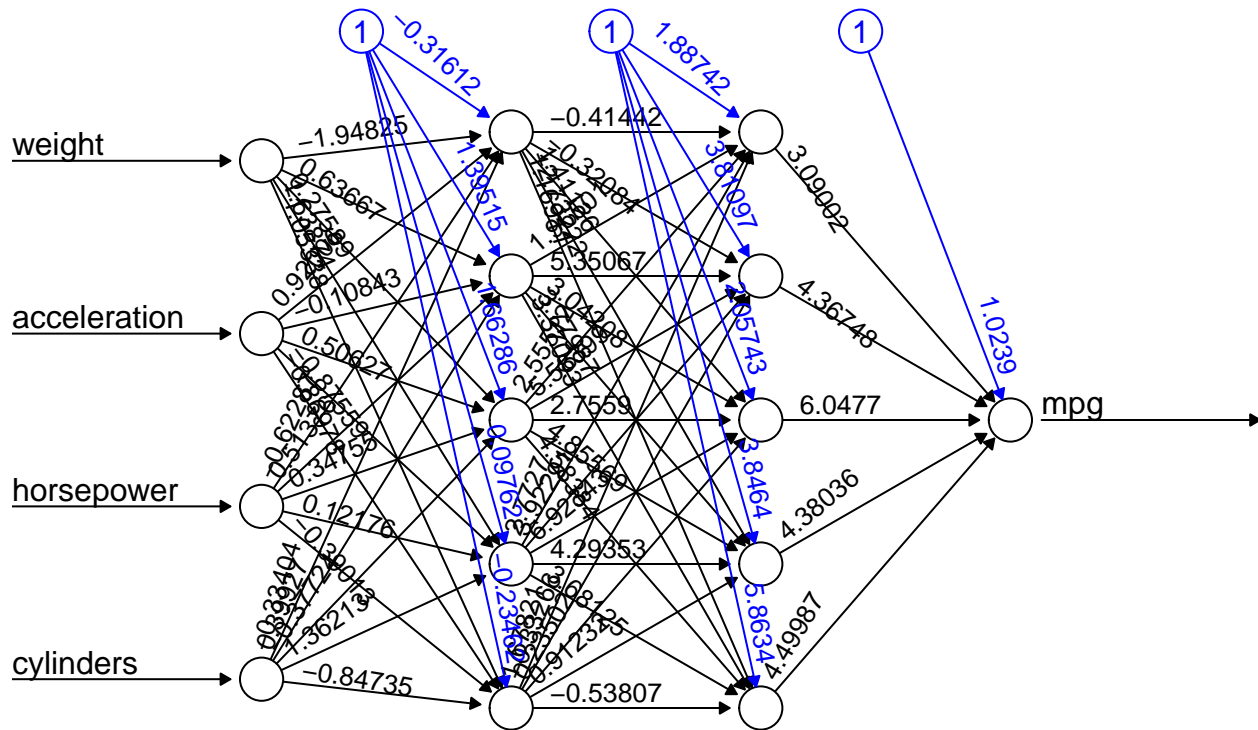
Error: 6050.2119 Steps: 252

or two hidden layers with 5 nodes each

```
nn55 =neuralnet(mpg ~ weight + acceleration + horsepower + cylinders, data=Auto.train, hidden=c(5,5))
summary(nn55)
```

```
##          Length Class      Mode
## call          4    -none-    call
## response     200    -none-   numeric
## covariate     800    -none-   numeric
## model.list      2    -none-    list
## err.fct        1    -none-   function
## act.fct        1    -none-   function
## linear.output   1    -none-   logical
## data           9  data.frame list
## net.result      1    -none-    list
## weights        1    -none-    list
## startweights    1    -none-    list
## generalized.weights 1    -none-    list
## result.matrix   64    -none-   numeric
```

```
#nn55
plot(nn55, rep="best")
```



Error: 6050.2119 Steps: 69

Which Neural Network has the best prediction?

```
Predict0 = compute( nn0, subset(Auto.test, select=c(weight, acceleration, horsepower, cylinders) ) )

# This prediction consists of X-variables "neurons" and predicted Y-variable "net.result".

names(Predict0)

## [1] "neurons"      "net.result"
mean( (Auto.test$mpg - Predict0$net.result)^2 )

## [1] 17.8324906
# Prediction MSE of this ANN is 18.83.

Predict5 = compute( nn5, subset(Auto.test, select=c(weight, acceleration, horsepower, cylinders) ) )
mean( (Auto.test$mpg - Predict5$net.result)^2 )

## [1] 61.03699646
# Its 3-node competitor has a much higher prediction MSE.

Predict55 = compute( nn55, subset(Auto.test, select=c(weight, acceleration, horsepower, cylinders) ) )
mean( (Auto.test$mpg - Predict55$net.result)^2 )

## [1] 61.03699901
```

This more complicated network should give us better results, not worse!! What is wrong?


Let's try scaling the data before we train our neural network.

```
AutoSc <- Auto
AutoSc$mpg <- scale(Auto$mpg, center = min(Auto$mpg), scale = max(Auto$mpg)-min(Auto$mpg))
AutoSc$weight <- scale(Auto$weight, center = min(Auto$weight), scale = max(Auto$weight)-min(Auto$weight))
AutoSc$acceleration <- scale(Auto$acceleration, center = min(Auto$acceleration), scale = max(Auto$acceleration)-min(Auto$acceleration))
AutoSc$horsepower <- scale(Auto$horsepower, center = min(Auto$horsepower), scale = max(Auto$horsepower)-min(Auto$horsepower))
AutoSc$cylinders <- scale(Auto$cylinders, center = min(Auto$cylinders), scale = max(Auto$cylinders)-min(Auto$cylinders))

trainData <- AutoSc[Z,]
testData <- AutoSc[-Z,]
nn0 = neuralnet(mpg ~ weight + acceleration + horsepower + cylinders, data=trainData, hidden=0)
summary(nn0)
```

| ## | Length | Class | Mode |
|------------------------|--------|------------|----------|
| ## call | 4 | -none- | call |
| ## response | 200 | -none- | numeric |
| ## covariate | 800 | -none- | numeric |
| ## model.list | 2 | -none- | list |
| ## err.fct | 1 | -none- | function |
| ## act.fct | 1 | -none- | function |
| ## linear.output | 1 | -none- | logical |
| ## data | 9 | data.frame | list |
| ## net.result | 1 | -none- | list |
| ## weights | 1 | -none- | list |
| ## startweights | 1 | -none- | list |
| ## generalized.weights | 1 | -none- | list |
| ## result.matrix | 8 | -none- | numeric |

```
plot(nn0, rep="best")
```



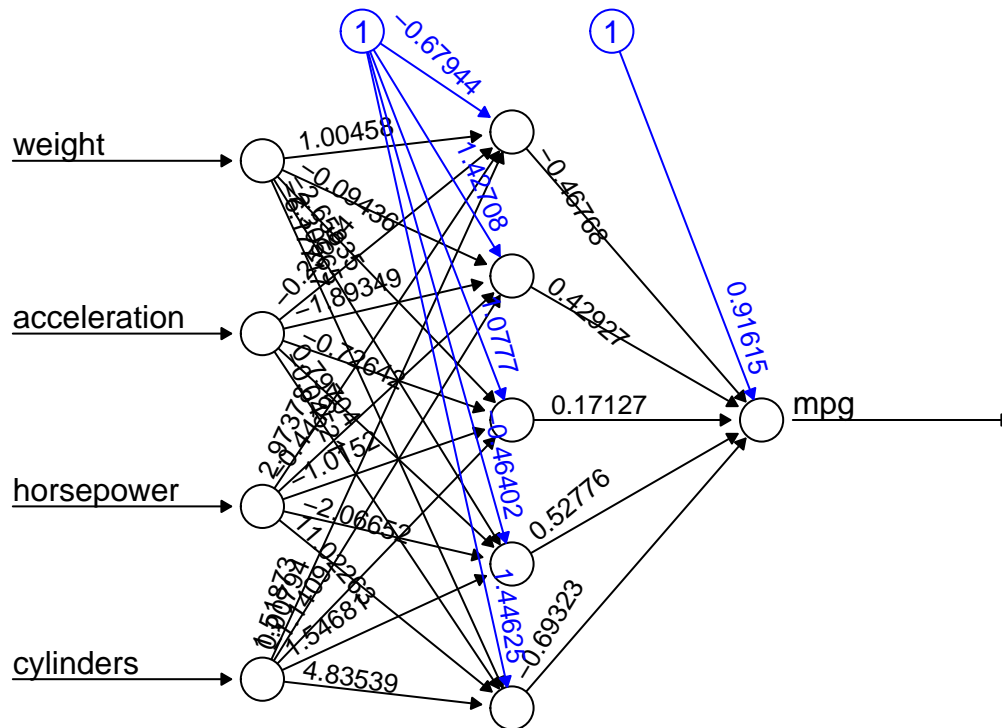
Error: 1.294473 Steps: 526

```
pr.nn0 = compute( nn0, subset(testData, select=c(weight, acceleration, horsepower, cylinders)) )
```

```
nn5 = neuralnet(mpg ~ weight + acceleration + horsepower + cylinders, data=trainData, hidden=5)
summary(nn5)
```

| ## | Length | Class | Mode |
|------------------------|--------|------------|----------|
| ## call | 4 | -none- | call |
| ## response | 200 | -none- | numeric |
| ## covariate | 800 | -none- | numeric |
| ## model.list | 2 | -none- | list |
| ## err.fct | 1 | -none- | function |
| ## act.fct | 1 | -none- | function |
| ## linear.output | 1 | -none- | logical |
| ## data | 9 | data.frame | list |
| ## net.result | 1 | -none- | list |
| ## weights | 1 | -none- | list |
| ## startweights | 1 | -none- | list |
| ## generalized.weights | 1 | -none- | list |
| ## result.matrix | 34 | -none- | numeric |

```
plot(nn5, rep="best")
```



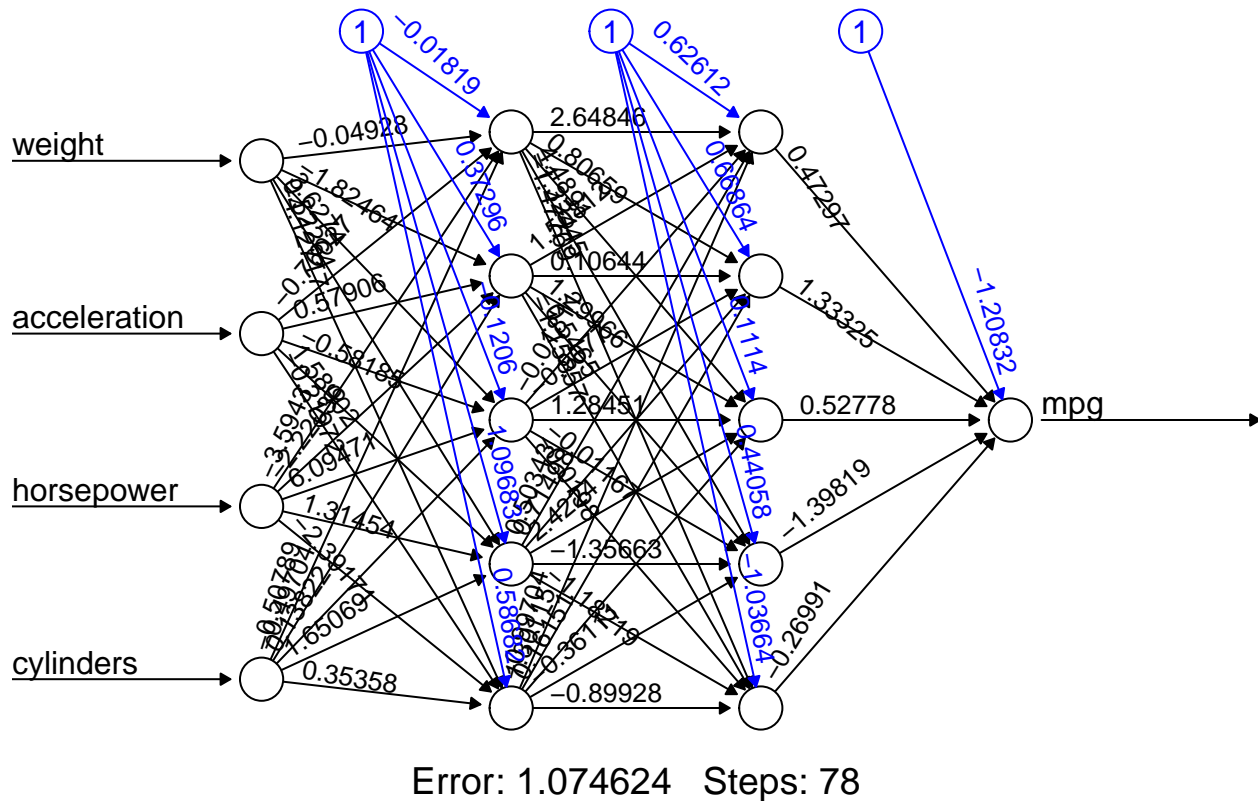
Error: 1.03703 Steps: 404

```
pr.nn5 = compute( nn5, subset(testData, select=c(weight, acceleration, horsepower, cylinders)) ) )

nn55 = neuralnet(mpg ~ weight + acceleration + horsepower + cylinders, data=trainData, hidden=c(5,5))
summary(nn55)
```

| ## | Length | Class | Mode |
|------------------------|--------|------------|----------|
| ## call | 4 | -none- | call |
| ## response | 200 | -none- | numeric |
| ## covariate | 800 | -none- | numeric |
| ## model.list | 2 | -none- | list |
| ## err.fct | 1 | -none- | function |
| ## act.fct | 1 | -none- | function |
| ## linear.output | 1 | -none- | logical |
| ## data | 9 | data.frame | list |
| ## net.result | 1 | -none- | list |
| ## weights | 1 | -none- | list |
| ## startweights | 1 | -none- | list |
| ## generalized.weights | 1 | -none- | list |
| ## result.matrix | 64 | -none- | numeric |

```
plot(nn55, rep="best")
```

```
pr.nn55 = compute( nn55, subset(testData, select=c(weight, acceleration, horsepower, cylinders) ) )
# Results from NN are normalized (scaled)
# Descaling for comparison
pr.nn0 <- pr.nn0$net.result*(max(Auto$mpg)-min(Auto$mpg))+min(Auto$mpg)
pr.nn5 <- pr.nn5$net.result*(max(Auto$mpg)-min(Auto$mpg))+min(Auto$mpg)
pr.nn55 <- pr.nn55$net.result*(max(Auto$mpg)-min(Auto$mpg))+min(Auto$mpg)
test.r <- (testData$mpg)*(max(Auto$mpg)-min(Auto$mpg))+min(Auto$mpg)

# Calculating MSE
MSE.nn <- sum((test.r - pr.nn0)^2)/nrow(testData)
MSE.nn5 <- sum((test.r - pr.nn5)^2)/nrow(testData)
MSE.nn55 <- sum((test.r - pr.nn55)^2)/nrow(testData)
# Compare the MSEs
print(paste(MSE.lm,MSE.nn,MSE.nn5, MSE.nn55))
```

```
## [1] "17.8324736443234 17.8620306670795 14.5163572372295 15.2882860115222"
```

We see how critical it is to scale our data before building a neural network!!

Using neural networks to do classification

```
AutoSc$ECO = ifelse( AutoSc$mpg > 0.5, "Economy", "Consuming" )
n = length(AutoSc$mpg)
AutoSc$ECO4 = rep("Economy",n)
AutoSc$ECO4[AutoSc$mpg < 0.75] = "Good"
```

```

AutoSc$ECO4[AutoSc$mpg < 0.5] = "OK"
AutoSc$ECO4[AutoSc$mpg < 0.25] = "Consuming"
trainData <- AutoSc[Z,]
testData <- AutoSc[-Z,]

```

First we will use one output variable. Train an artificial neural network to classify cars into “Economy” and “Consuming”.

```

library(nnet)
nn.class = nnet( as.factor(ECO) ~ weight + acceleration + horsepower + cylinders, data=trainData, size=

```

```

## # weights: 31
## initial value 136.669815
## iter 10 value 61.420327
## iter 20 value 53.619878
## iter 30 value 52.161995
## iter 40 value 50.549721
## iter 50 value 47.125765
## iter 60 value 45.048199
## iter 70 value 44.533326
## iter 80 value 43.177338
## iter 90 value 40.768818
## iter 100 value 37.656840
## iter 110 value 36.883535
## iter 120 value 35.995247
## iter 130 value 35.623704
## iter 140 value 35.372456
## iter 150 value 35.030882
## iter 160 value 34.723306
## iter 170 value 33.544330
## iter 180 value 32.050584
## iter 190 value 31.212941
## iter 200 value 30.248805
## iter 210 value 29.751090
## iter 220 value 29.277282
## iter 230 value 28.562061
## iter 240 value 28.081954
## iter 250 value 27.582901
## iter 260 value 27.188632
## iter 270 value 27.010850
## iter 280 value 26.899537
## iter 290 value 26.726831
## iter 300 value 26.515571
## iter 310 value 26.163533
## iter 320 value 26.013003
## iter 330 value 25.938772
## iter 340 value 25.785247
## iter 350 value 25.733330
## iter 360 value 25.657508
## iter 370 value 25.448427
## iter 380 value 25.189268
## iter 390 value 25.068564
## iter 400 value 24.823322
## iter 410 value 24.675024

```

```
## iter 420 value 24.573237
## iter 430 value 24.279377
## iter 440 value 23.837118
## iter 450 value 23.776020
## iter 460 value 23.451186
## iter 470 value 23.145696
## iter 480 value 23.017139
## iter 490 value 22.837513
## iter 500 value 22.732903
## final value 22.732903
## stopped after 500 iterations
```

```
summary(nn.class)
```

```
## a 4-5-1 network with 31 weights
## options were - entropy fitting
## b->h1 i1->h1 i2->h1 i3->h1 i4->h1
## -1.97 -241.81 110.44 -104.02 107.53
## b->h2 i1->h2 i2->h2 i3->h2 i4->h2
## -49.78 -6.22 43.05 102.38 20.13
## b->h3 i1->h3 i2->h3 i3->h3 i4->h3
## 16.12 -73.86 -18.29 -26.72 13.78
## b->h4 i1->h4 i2->h4 i3->h4 i4->h4
## -84.48 9.81 73.13 163.75 22.20
## b->h5 i1->h5 i2->h5 i3->h5 i4->h5
## 60.38 -191.88 -52.98 -117.51 -12.11
## b->o h1->o h2->o h3->o h4->o h5->o
## -3.12 4.81 -77.67 133.58 74.30 -92.33
```

```
pr.nn = predict(nn.class, subset(testData, select=c(weight,acceleration,horsepower,cylinders) ), type =
head(pr.nn)
```

```
## [1] "Consuming" "Consuming" "Consuming" "Consuming" "Consuming" "Consuming"
```

```
table(pr.nn,testData$ECO)
```

```
##
## pr.nn      Consuming Economy
## Consuming      112      18
## Economy        21      41
```

```
# train neural network on 4 inputs, 1 hidden layer with 5 nodes, 4 output nodes
```

```
nn4.class = nnet( as.factor(ECO4) ~ weight+acceleration+horsepower+cylinders, data=trainData, size=5 , n
```

```
## # weights: 49
## initial value 285.788088
## iter 10 value 147.621204
## iter 20 value 127.164456
## iter 30 value 119.910773
## iter 40 value 110.919763
## iter 50 value 105.414131
## iter 60 value 102.323318
## iter 70 value 100.340219
## iter 80 value 98.912595
## iter 90 value 97.952976
## iter 100 value 96.803550
## iter 110 value 95.867543
```

```

## iter 120 value 94.644272
## iter 130 value 93.560171
## iter 140 value 92.840335
## iter 150 value 92.468186
## iter 160 value 92.260835
## iter 170 value 92.137376
## iter 180 value 91.895166
## iter 190 value 91.740578
## iter 200 value 91.584289
## iter 210 value 91.567647
## iter 220 value 91.521571
## iter 230 value 91.485273
## iter 240 value 91.424759
## iter 250 value 91.395310
## iter 260 value 91.364708
## iter 270 value 91.351189
## iter 280 value 91.309009
## iter 290 value 91.232543
## iter 300 value 91.183813
## iter 310 value 91.175779
## iter 320 value 91.156833
## iter 330 value 91.146345
## iter 340 value 91.136255
## iter 350 value 91.133364
## iter 360 value 91.129454
## iter 370 value 91.123012
## iter 380 value 91.110374
## iter 390 value 91.097234
## iter 400 value 91.092726
## iter 410 value 91.092145
## iter 420 value 91.079176
## iter 430 value 91.073899
## iter 440 value 91.071119
## iter 450 value 91.069591
## iter 460 value 91.068497
## iter 470 value 91.063148
## iter 480 value 91.060869
## iter 490 value 91.056823
## iter 500 value 91.027115
## final value 91.027115
## stopped after 500 iterations

```

```
summary(nn4.class)
```

```

## a 4-5-4 network with 49 weights
## options were - softmax modelling
## b->h1 i1->h1 i2->h1 i3->h1 i4->h1
## -3.71 13.80 1.74 -2.04 -3.85
## b->h2 i1->h2 i2->h2 i3->h2 i4->h2
## 3.43 -14.59 -1.89 6.41 2.05
## b->h3 i1->h3 i2->h3 i3->h3 i4->h3
## -2.97 8.59 0.63 3.87 -2.94
## b->h4 i1->h4 i2->h4 i3->h4 i4->h4
## 75.79 88.60 -129.59 -8.30 24.73
## b->h5 i1->h5 i2->h5 i3->h5 i4->h5

```

```
## -17.67    5.47    1.33   -8.25   30.30
##  b->o1  h1->o1  h2->o1  h3->o1  h4->o1  h5->o1
## -32.21    8.14   24.21   28.65    1.47   19.96
##  b->o2  h1->o2  h2->o2  h3->o2  h4->o2  h5->o2
##  70.45  -74.09  -73.26  -24.71    5.49    8.58
##  b->o3  h1->o3  h2->o3  h3->o3  h4->o3  h5->o3
## -53.75   97.66   61.93  -38.24   -1.10  -37.17
##  b->o4  h1->o4  h2->o4  h3->o4  h4->o4  h5->o4
##  15.84  -32.93  -13.61   35.39   -5.85    9.39
```

```
names(nn4.class)
```

```
## [1] "n"           "nunits"      "nconn"       "conn"
## [5] "nsunits"     "decay"       "entropy"     "softmax"
## [9] "censored"    "value"       "wts"         "convergence"
## [13] "fitted.values" "residuals"   "lev"         "call"
## [17] "terms"       "coefnames"   "xlevels"
```

```
pr.nn4 = predict(nn4.class, subset(testData, select=c(weight,acceleration,horsepower,cylinders) ), type="class")
head(pr.nn4)
```

```
## [1] "OK"           "Consuming" "Consuming" "Consuming" "Consuming" "Consuming"
```

```
table(pr.nn4,testData$EC04)
```

```
##
## pr.nn4      Consuming Economy Good OK
## Consuming      53         0    0  9
## Economy         0         3    3  3
## Good            0         6   26 16
## OK              9         0   21 43
```