

KNN – K-nearest neighbor classification

Data – miles per gallon and other variables from the Auto data set.

```
> library(ISLR) # This library contains datasets from our textbook (ISLR = name of our text)
> attach(Auto)
> names(Auto) # List of variables in this dataset
[1] "mpg" "cylinders" "displacement" "horsepower" "weight" "acceleration" "year" "origin" "name"
> summary(mpg) # Economy rating will be defined based on miles per gallon
Min. 1st Qu. Median Mean 3rd Qu. Max.
9.00 17.00 22.75 23.45 29.00 46.60
# Initiate a fuel consumption rating variable that will be treated as categorical
> Economy = rep("Gas consumption", length(mpg))
> Economy[mpg <= 17] = "Heavy"
> Economy[mpg > 17 & mpg <= 22.75] = "OK"
> Economy[mpg > 22.75 & mpg <= 29] = "Eco"
> Economy[mpg > 29] = "Excellent"

> table(Economy) # We used sample quartiles of variable mpg to define these ratings,
Economy # that's why we got four approximately equal groups.
# Now, we'll derive a classification rule, using other car characteristics
Eco Excellent Heavy OK
101 95 99 97
```

Prepare training and testing data, predictors (X) and responses (Y)

```
> n = length(mpg)
> Z = sample(n, n/2) # We'll split data at random

> Auto.training = Auto[Z, ] # Subsample with indices from that subsample Z
> Auto.testing = Auto[-Z, ] # Notice the "minus" sign to denote all indices except those in Z
> dim(Auto)
[1] 392 9
> dim(Auto.training)
[1] 196 9
> dim(Auto.testing)
[1] 196 9

> names(Auto)
[1] "mpg" "cylinders" "displacement" "horsepower" "weight" "acceleration" "year" "origin" "name"
```

KNN in R requires 4 inputs: training X, testing X, training Y, and K.

```
> X.training = Auto.training[, 2:7] # Take columns (variables) 2-7. That's from cylinders to year.
> X.testing = Auto.testing[, 2:7]
> Y.training = Economy[Z]
> Y.testing = Economy[-Z]
```

KNN tool is in the package "class".

```
> library(class)
> knn.result = knn( X.training, X.testing, Y.training, 3 )
```

```
> table( Y.testing, knn.result )
```

Y.testing \ knn.result	Eco	Excellent	Heavy	OK
Eco	19	17	1	11
Excellent	9	35	0	2
Heavy	0	0	38	8
OK	5	3	5	32

```
> mean( Y.testing == knn.result )
```

```
[1] 0.6702703 # 67% correct classification rate with K=3. Is there a better K?
```

```
# We'll check all K from 1 to 20.
```

```
> class.rate = rep(0,20) # Create a vector of length 20 and fill it with classification rates,
# computed in a do-loop
```

```
> for (K in 1:20) {
+ knn.result = knn( X.training, X.testing, Y.training, K )
+ class.rate[K]=mean( Y.testing == knn.result )
+ }
```

```
> class.rate
```

Apparently, K=6 and K=8 provide a slightly better prediction although still not as good as LDA

```
[1] 0.6378378 0.6378378 0.6702703 0.6810811 0.6810811 0.6918919 0.6702703
[8] 0.6918919 0.6648649 0.6648649 0.6594595 0.6594595 0.6756757 0.6702703
[15] 0.6702703 0.6864865 0.6702703 0.6702703 0.6702703 0.6702703
```