


## **# CROSS-VALIDATION**

### **1. Validation Set Approach**

```
> library(ISLR2)
> attach(Auto); n = length(mpg);
> n
[1] 392
> Z = sample( n, 250 )           # Random subsample of size 250 (does not have to be
n/2)                             # Works similarly to generating a Bernoulli variable

> reg.fit = lm( mpg ~ weight + horsepower + acceleration, subset=Z )      # Fit using training data
> mpg_predicted = predict( reg.fit, newdata=Auto[-Z, ] )                  # Use this model to predict the testing data [-Z]
                                                                            # We can see a nonlinear component
> plot(mpg[-Z],mpg_predicted)                                             # Compare with the line y=x.
> abline(0,1)
```



```
> mean( (mpg[-Z] - mpg_predicted)^2 )           # Estimate the mean-squared error MSE
[1] 18.58128
```

### **2. Jackknife (Leave-One-Out Cross-Validation, LOOCV)**

**“Manually”:**

```
> Yhat = numeric(n)
> for (i in 1:n){
+ reg = lm( mpg ~ weight + horsepower + acceleration, data=Auto[-i,] )
+ Yhat[i] = predict( reg, newdata=Auto[i,] )
+ }
> plot(mpg,Yhat)
```

```
> mean((mpg-Yhat)^2)
[1] 18.25595
```

Using package “boot”:


```
> install.packages("boot")
> library(boot)
> glm.fit = glm(mpg ~ weight + horsepower + acceleration) # Default “family” is Normal, so this GLM model
> cv.error = cv.glm( Auto, glm.fit ) # is the same as LM, standard linear regression.
> names(cv.error) # This cross-validation tool has several outputs
[1] "call" "K" "delta" "seed" # We are interested in “delta”

> error$delta
[1] 18.25595 18.25542
```

# Delta consists of 2 numbers – estimated prediction error and its version adjusted for the lost sample size due to cross-validation.

# Example – what power of “horsepower” is optimal for this prediction?

```
> cv.error = rep(0,10) # Initiate a vector of estimated errors
> for (p in 1:10){ # Fit polynomial regression models
+ glm.fit = glm( mpg ~ weight + poly(horsepower,p) + acceleration ) # with power p=1..10
+ cv.error[p] = cv.glm( Auto, glm.fit )$delta[1] } # Save prediction errors
> cv.error # Look at the results
[1] 18.25595 15.90163 15.72995 15.86879 15.74517 15.74989 15.70073 15.79314 15.95933 16.38301
# Although p=7 yields the lowest estimated prediction error, after p=2, the improvement is very little.
> plot(cv.error)
> lines(cv.error)
```



### 3. K-Fold Cross-Validation

**# We can specify K within cv.glm (Omitted K is K=1 by default, which is LOOCV).**

```
> cv.error = rep(0,10)
> for (p in 1:10){ glm.fit = glm( mpg ~ weight + poly(horsepower,p) + acceleration )
+ cv.error[p] = cv.glm( Auto, glm.fit, K=30 )$delta[1] }
> cv.error
[1] 18.22699 15.87030 15.73565 15.85911 15.80686 15.71585 16.09625 15.67797 15.82097 16.48196
> which.min(cv.error)
[1] 6
```

#### **4. Cross-validation in classification problems. Loss function.**

**# Command cv.glm, as we used it above, calculates MSE, the mean squared error, and calls them “delta”.**  
**# In classification problems, MSE can be used to measure the distance between the response variable Y**  
**# and the predicted probability p. For this, Y has to be 0 or 1, and p should be the probability of Y=1.**  
**# However, the correct classification rate and the error classification rate are more standard measures**  
**# of classification accuracy. We can force cv.glm to return these measures by introducing a suitable loss**  
**# function. For example, we’ll be predicting whether a student is depressed or not.**

**# We define a loss function  $L(Y,p)$ , a function of true response Y and predicted probability p. The loss = 1 if**  
**# the predicted response is different from the actual response and 0 otherwise. Suppose the threshold is 0.5**

```
> loss = function(Y,p){ return( mean( (Y==1 & p < 0.5) | (Y==0 & p >= 0.5) ) ) }
> loss(1,0.3)
[1] 1
> loss(c(1,1),c(0.3,0.7))
[1] 0.5
```

**# Now we attach the Depression data, skip missing values, fit logistic regression model, and estimate**  
**# the error classification rate by LOOCV.**

```
> Depr = read.csv(url("http://fs2.american.edu/~baron/627/R/depression_data.csv"))
> D = na.omit(Depr)
> attach(D)
> library(boot)
> lreg = glm(Diagnosis ~ Gender + Guardian_status + Cohesion_score, family="binomial")
> cv = cv.glm( D, lreg, loss )
> cv$delta[1]
[1] 0.1615721
```

**# Is Guardian\_status significant? Let’s compare the error rates with and without variable “Guardian\_status”.**

```
> lreg = glm(Diagnosis ~ Gender + Cohesion_score, family="binomial")
> cv.glm( D, lreg, loss )$delta[1]
[1] 0.1572052
```

**# The error classification rate is lower without the “Guardian\_status”.**