

DIMENSION REDUCTION AND SHRINKAGE

Part I. Variable Selection and Ridge Regression

1. VARIABLE SELECTION

```
> attach(Auto)
> library(leaps)
> reg.fit = regsubsets( mpg ~ cylinders + displacement + horsepower + weight + acceleration + year, Auto )
> summary(reg.fit)
Selection Algorithm: exhaustive
      cylinders displacement horsepower weight acceleration year
1 ( 1 ) " " " " " " "*" " "
2 ( 1 ) " " " " " " "*" " "
3 ( 1 ) " " " " " " "*" "*"
4 ( 1 ) " " "*" " " "*" "*"
5 ( 1 ) "*" "*" " " "*" "*"
6 ( 1 ) "*" "*" "*" "*" "*" "
```

This command finds the best model for each p = number of independent variables. The best model is determined by the lowest RSS.

Next, choose the best p according to some criteria:

```
> summary(reg.fit)$adjr2      # Adjusted  $R^2$ 
[1] 0.6918423 0.8071941 0.8071393 0.8067872 0.8067841 0.8062826
> summary(reg.fit)$cp        # Mallows  $C_p$ 
[1] 232.396144 1.169751 2.284200 3.992019 5.000800 7.000000
> summary(reg.fit)$bic        # BIC = Bayesian information criterion
[1] -450.5016 -629.3564 -624.2828 -618.6081 -613.6448 -607.6743
```

Recall that plain R^2 is not a fair measure of performance. It always increases with p :



```
> summary(reg.fit)$rsq
[1] 0.6926304 0.8081803 0.8086190 0.8087638 0.8092549 0.8092553
```

For stepwise or backward elimination variable selection, use method="forward" or method="backward".

```
> library(MASS)
> reg = regsubsets( medv ~ ., data=Boston, method = "backward" )
```


There is a nice way to visualize results, ranking models by the chosen "scale". Black color means the variable is included into the model, white means it is excluded.

```
> plot(reg)
> plot(reg, scale = "adjr2" )
```



To see more models, use option "nbest", which is the number of models of each size p to be compared.

```
> reg = regsubsets( medv ~ ., data=Boston, method = "backward", nbest=4 )
> plot(reg, scale = "adjr2" )
```



We can also choose the best model by means of a stepwise procedure, starting with one model and ending with another.

```
> null = lm( medv ~ 1, data=Boston )
> full = lm( medv ~ ., data=Boston )
> step( null, scope=list(lower=null, upper=full), direction="forward" )
Start: AIC=2246.51
medv ~ 1
```

	Df	Sum of Sq	RSS	AIC
+ lstat	1	23243.9	19472	1851.0
+ rm	1	20654.4	22062	1914.2
+ ptratio	1	11014.3	31702	2097.6
+ indus	1	9995.2	32721	2113.6
+ tax	1	9377.3	33339	2123.1
+ nox	1	7800.1	34916	2146.5
+ crim	1	6440.8	36276	2165.8
+ rad	1	6221.1	36495	2168.9
+ age	1	6069.8	36647	2171.0
+ zn	1	5549.7	37167	2178.1
+ black	1	4749.9	37966	2188.9
+ dis	1	2668.2	40048	2215.9
+ chas	1	1312.1	41404	2232.7
<none>			42716	2246.5

Compare contributions of
remaining independent variables

Step: AIC=1851.01
medv ~ lstat

	Df	Sum of Sq	RSS	AIC
+ rm	1	4033.1	15439	1735.6
+ ptratio	1	2670.1	16802	1778.4
+ chas	1	786.3	18686	1832.2
+ dis	1	772.4	18700	1832.5
+ age	1	304.3	19168	1845.0
+ tax	1	274.4	19198	1845.8
+ black	1	198.3	19274	1847.8
+ zn	1	160.3	19312	1848.8
+ crim	1	146.9	19325	1849.2
+ indus	1	98.7	19374	1850.4
<none>			19472	1851.0
+ rad	1	25.1	19447	1852.4
+ nox	1	4.8	19468	1852.9

... < truncated > ...

Step: AIC=1585.76
medv ~ lstat + rm + ptratio + dis + nox + chas + black + zn +
 crim + rad + tax

	Df	Sum of Sq	RSS	AIC
<none>			11081	1585.8
+ indus	1	2.51754	11079	1587.7
+ age	1	0.06271	11081	1587.8

Call:
lm(formula = medv ~ lstat + rm + ptratio + dis + nox + chas +
 black + zn + crim + rad + tax, data = Boston)

Coefficients:

(Intercept)	lstat	rm	ptratio	dis	nox
36.341145	-0.522553	3.801579	-0.946525	-1.492711	-17.376023
chas	black	zn	crim	rad	tax
2.718716	0.009291	0.045845	-0.108413	0.299608	-0.011778

The final model contains variables lstat, rm, ptratio, dis, nox, chas, black, zn, crim, rad, and tax.

2. RIDGE REGRESSION

Dataset “Boston” has some strong correlations, resulting in multicollinearity.

> cor(Boston)


Apply ridge regression

```
> lm.ridge( medv~., data=Boston, lambda=0.5 )
```

	crim	zn	indus	chas
36.270091954	-0.107524388	0.046092635	0.018815242	2.693518778
	nox	rm	age	dis
-17.646013524	3.816019080	0.000578687	-1.469191794	0.301928336
	tax	ptratio	black	lstat
-0.012134702	-0.950831885	0.009309388	-0.523845421	

We can see how the slopes change with penalty lambda. These are estimated slopes for different lambda. When lambda=0, we get LSE. Large lambda forces them toward 0.

```
> rr = lm.ridge(medv ~ ., data=train, lambda=seq(0,1000,1) )
> rr
> plot(rr)
```



To choose a good lambda, fit ridge regression with various lambda and compare prediction performance.

```
> select(rr)
modified HKB estimator is 4.321483
modified L-W estimator is 3.682107
smallest value of GCV at 5
```

So, the best lambda is around 5. We'll look closer at that range.

```
> rr = lm.ridge( medv~., data=Boston, lambda=seq(0,10,0.01) )
```

```
> select(rr)
modified HKB estimator is 4.594163
modified L-W estimator is 3.961575
smallest value of GCV at 4.26
> plot(rr$lambda,rr$GCV)
```

