

DIMENSION REDUCTION AND SHRINKAGE

Part II. Ridge Regression and LASSO

3. LASSO and RIDGE REGRESSION in package GLMNET

```
> library(glmnet)
```

This package requires X-variables in a matrix

```
> X = model.matrix( medv ~ ., data=Boston )  
> Y = medv  
> ridgereg = glmnet(X, Y, alpha=0, lambda = seq(0,10,0.01))
```

alpha is a “mixing parameter”. It combines Lasso and Ridge Regression. We only need the extreme values for now,

| | | |
|---------|----|------------------|
| alpha=0 | => | ridge regression |
| alpha=1 | => | lasso |

So, which lambda is it best to choose? Run cross-validation...


```
> cv_ridge = cv.glmnet(X,medv,alpha=0,lambda=seq(0,10,0.01))  
> names(cv_ridge)  
[1] "lambda" "cvm" "cvstd" "cvup" "cvlo" "nzero" "name" "glmnet.fit"  
[9] "lambda.min" "lambda.1se"
```

For the selected values of lambda, we get

- “cvm” = the mean cross-validation error
- “cvstd” = its estimated standard error
- “cvlo” = cvm – cvstd (lower curve)
- “cvup” = cvm + cvstd (upper curve)

All these can be plotted...

```
> plot(cv_ridge)
```



Which lambda minimized the MSE?

```
> cv_rideg$lambda.min
[1] 0.12
> predict( ridgereg, cv_rideg$lambda.min, type="coefficients" )
```



```
(Intercept) 20.863586526
crim        -0.068099108
zn          0.020536577
indus       -0.068164404
chas        2.777651593
nox         -5.514335229
rm          3.645249869
age         -0.007642606
dis         -0.532994082
rad         0.034679405
tax         -0.002916331
ptratio     -0.676296397
black       0.007674913
lstat      -0.351786530
```

Similarly with LASSO, only choose alpha=1

```
> lasso = glmnet(X, Y, alpha=1, lambda = seq(0,10,0.01))
```


Compare the slopes estimated by ridge regression and by lasso

```
> plot(ridgereg)
> plot(lasso)
```



Ridge regression uses all the variables – all slopes are not 0 for all lambda. Conversely, lasso does variable selection and sends some slopes to 0. The number of non-zero slopes is printed in the top.

```
> cv.lasso = cv.glmnet( X, medv, alpha=1, lambda=seq(0,10,0.01) )
> plot(cv.lasso)
```



```

> cv.lasso$lambda.min
[1] 0.02
> predict( lasso, cv. lasso$lambda.min, type="coefficients" )
(Intercept) 18.739101467
crim        -0.024356546
zn          .
indus       .
chas        2.009577446
nox         -4.667589527
rm          4.273554725
age         .
dis         -0.401567952
rad         .
tax         .
ptratio     -0.803881292
black       0.006716721
lstat       -0.518576315

```

For LASSO, the best lambda to use is 0.02. Some coefficients are 0 – these variables are removed from the model.

Prediction for new values of X and cross-validation

```

> n = length(medv)
> Z = sample(n,n/2)

> lasso = glmnet( X[Z,], medv[Z], alpha=1, lambda=seq(0,10,0.01) )
> Yhat = predict( lasso, cv.lasso$lambda.min, newx=X[-Z,] )
> mean((Yhat - medv[-Z])^2)
[1] 23.62894

```

This is the test MSE, estimated by the validation set approach.