

SUPPORT VECTOR MACHINES


1. SVM with various kernels

The SVM command is in package called e1071.

```
> install.packages("e1071");  
> library(e1071)
```

Let's use support vector machines to classify cars into Economy and Consuming classes.

```
> ECO = ifelse( mpg > 22.75, "Economy", "Consuming" )  
> Color = ifelse( mpg > 22.75, "green", "red" )  
> plot( weight, horsepower, lwd=3, col=Color )
```



The two classes cannot be separated by a hyperplane, but the SVM method is surely applicable.

```
> S = svm( ECO ~ weight + horsepower, data=Auto, kernel = "linear" )  
Error in svm.default(x, y, scale = scale, ..., na.action = na.action) :  
  Need numeric dependent variable for regression.
```

Error? There are other, unused variables in dataset Auto that prevent R from doing this SVM analysis. We'll create a reduced dataset.

```
> d = data.frame(ECO, weight, horsepower)  
  
> S = svm( ECO ~ weight + horsepower, data=d, kernel="linear" )  
> summary(S)
```

Parameters:

SVM-Type: C-classification

SVM-Kernel: linear

cost: 1

gamma: 0.5

Number of Support Vectors: 120

(60 60)


So, there are 120 points violating the separating hyperplane or the margin, 60 in each class.

```
> plot(S, data=Auto)
```

Error in plot.svm(S, data = Auto) : missing formula.

Same story. We need to use a reduced dataset that contains only the needed variables.

```
> plot(S, data=d)
```



This is the final classification with a linear kernel and therefore, a linear boundary. Support vectors are marked as "x", other points as "o".

We can look at other types of kernels and boundaries – polynomial, radial, and sigmoid.

```
> S = svm( ECO ~ weight + horsepower, data=d, kernel="polynomial" )
```

```
> summary(S); plot(S,d)
```

Number of Support Vectors: 176

```
> S = svm( ECO ~ weight + horsepower, data=d, kernel="radial" )
```

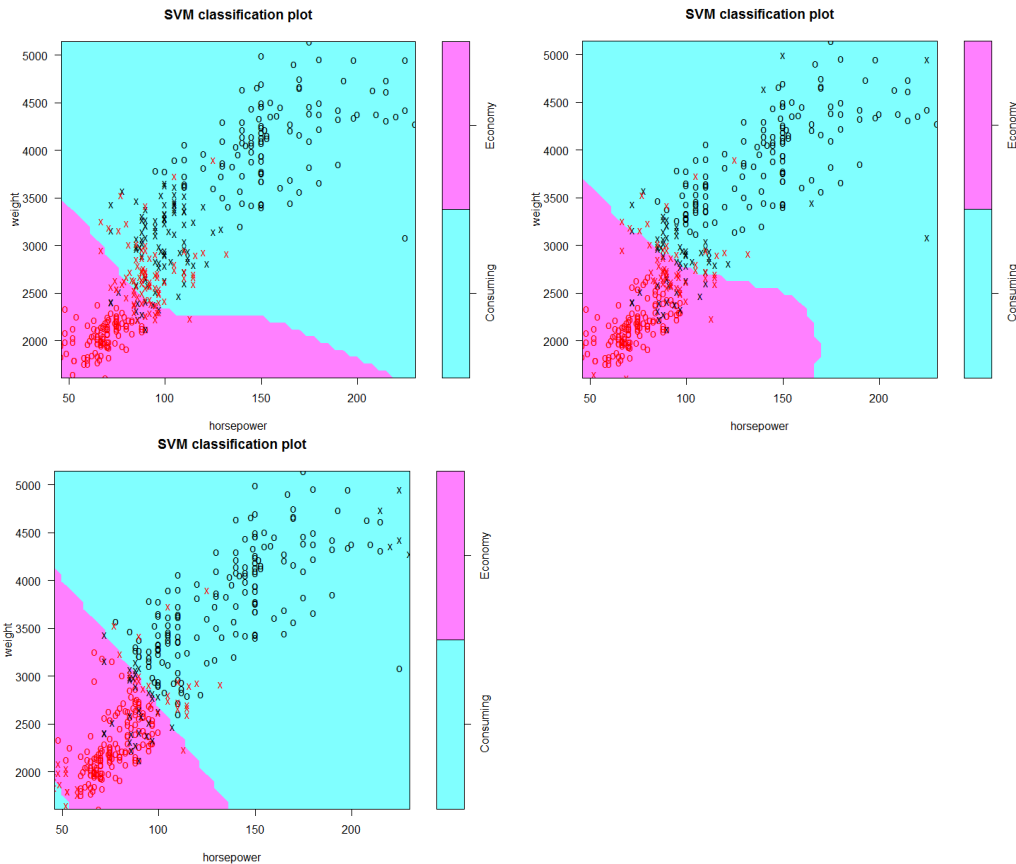
```
> summary(S); plot(S,d)
```

Number of Support Vectors: 121

```
> S = svm( ECO ~ weight + horsepower, data=d, kernel="sigmoid" )
```

```
> summary(S); plot(S,d)
```

Number of Support Vectors: 74



Adding more variables should give a better fit – to the *training* data.

```
> S = svm( factor(ECO) ~ weight + horsepower + displacement + cylinders, data=Auto, kernel="linear" )
```

```
> summary(S)
```

Number of Support Vectors: 99

We can identify the support vectors:

```
> S$index
```

```
[1] 16 17 18 25 33 45 46 48 60 61 71 76 77 78 80 100 107 108 109  
[20] 110 111 112 113 119 120 123 153 154 162 173 178 199 206 208 209 210 240 241  
[39] 242 253 258 262 269 273 274 275 280 281 384 24 31 49 84 101 114 122 131  
[58] 149 170 177 179 192 205 218 233 266 270 271 296 297 298 299 305 306 313 314  
[77] 318 322 326 327 331 337 338 353 355 356 357 358 360 363 365 368 369 375 381  
[96] 382 383 385 387
```

```
> Auto[S$index,]
```

mpg cylinders displacement horsepower weight acceleration year origin

16	22.0	6	198	95	2833	15.5	70	1
17	18.0	6	199	97	2774	15.5	70	1
18	21.0	6	200	85	2587	16.0	70	1
25	21.0	6	199	90	2648	15.0	70	1

< truncated >

2. Tuning and cross-validation

The “cost” option specifies the cost of violating the margin. We can try costs 0.001, 0.01, 0.1, 1, 10, 100, 1000:

```
> Stuned = tune( svm, ECO ~ weight + horsepower, data=d, kernel="linear", ranges=list(cost=10^seq(-3,3)) )
> summary(Stuned)
- sampling method: 10-fold cross validation
```

- best parameters:

```
cost
0.1
```

- best performance: 0.1173718

- Detailed performance results:

```
cost error dispersion
1 1e-03 0.2478205 0.10663023
2 1e-02 0.1432051 0.05485355
3 1e-01 0.1173718 0.04208311
4 1e+00 0.1326282 0.04461101
5 1e+01 0.1351923 0.04819639
6 1e+02 0.1351923 0.04819639
7 1e+03 0.1351923 0.04819639
```

This cost yielded the lowest cross-validation error of

classification.

We can also find the optimal kernel.

```
> Stuned = tune( svm, ECO ~ weight + horsepower, data=d, ranges=list(cost=10^seq(-3,3),
kernel=c("linear","polynomial","radial","sigmoid"))) )
> summary(Stuned)
```

Parameter tuning of ‘svm’:

- sampling method: 10-fold cross validation

- best parameters:

```
cost kernel
0.1 sigmoid
```

- best performance: 0.1046154

- Detailed performance results:

```
cost kernel error dispersion
1 1e-03 linear 0.2164744 0.10501351
2 1e-02 linear 0.1326282 0.05074006
```

3	1e-01	linear	0.1096154	0.04330918
4	1e+00	linear	0.1172436	0.03813782
5	1e+01	linear	0.1223718	0.04775672
6	1e+02	linear	0.1223718	0.04775672
7	1e+03	linear	0.1223718	0.04775672
8	1e-03	polynomial	0.3720513	0.08274072
9	1e-02	polynomial	0.2601282	0.06438244
10	1e-01	polynomial	0.1987821	0.07443903
11	1e+00	polynomial	0.1784615	0.05328633
12	1e+01	polynomial	0.1580769	0.04909157
13	1e+02	polynomial	0.1555128	0.04999836
14	1e+03	polynomial	0.1504487	0.04722372
15	1e-03	radial	0.5816026	0.05687780
16	1e-02	radial	0.1301282	0.05190241
17	1e-01	radial	0.1198077	0.05104329
18	1e+00	radial	0.1223718	0.04118608
19	1e+01	radial	0.1096795	0.04835338
20	1e+02	radial	0.1198718	0.04184981
21	1e+03	radial	0.1146795	0.04354410
22	1e-03	sigmoid	0.5816026	0.05687780
23	1e-02	sigmoid	0.1530769	0.04517581
24	1e-01	sigmoid	0.1046154	0.03711533
25	1e+00	sigmoid	0.1173718	0.04715638
26	1e+01	sigmoid	0.1530769	0.06159616
27	1e+02	sigmoid	0.1582051	0.06489946
28	1e+03	sigmoid	0.1582051	0.06489946

The best kernel and cost.

```
> Soptimal = svm( ECO ~ weight + horsepower, data=d, cost=0.1, kernel="sigmoid" )
```

```
> summary(Soptimal); plot(Soptimal,data=d)
```

Parameters:

SVM-Type: C-classification

SVM-Kernel: sigmoid

cost: 0.1

gamma: 0.5


Number of Support Vectors: 164
(82 82)

We know that more support vectors imply a lower variance

Number of Classes: 2

Levels: Consuming Economy

SVM classification plot



Let's use the validation set method to estimate the classification rate of this optimal SVM.

```
> n = length(mpg); Z = sample(n,n/2)
> Strain = svm( ECO ~ weight + horsepower, data=d[Z,], cost=0.1, kernel="sigmoid" )
> Yhat = predict( Strain, data=d[-Z,] )
> table( Yhat, ECO[Z] )
```

Yhat	Consuming	Economy
Consuming	82	9
Economy	17	88

```
> table( Yhat, ECO[Z] )
> mean( Yhat==ECO[Z] )
[1] 0.8673469
```

3. More than two classes

Let's create more categories of ECO. The same tool svm() can handle multiple classes.

```
> summary(mpg)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
   9.00  17.00  22.75  23.45  29.00  46.60
```

```
> ECO4 = rep("Economy",n)
> ECO4[mpg < 29] = "Good"
```

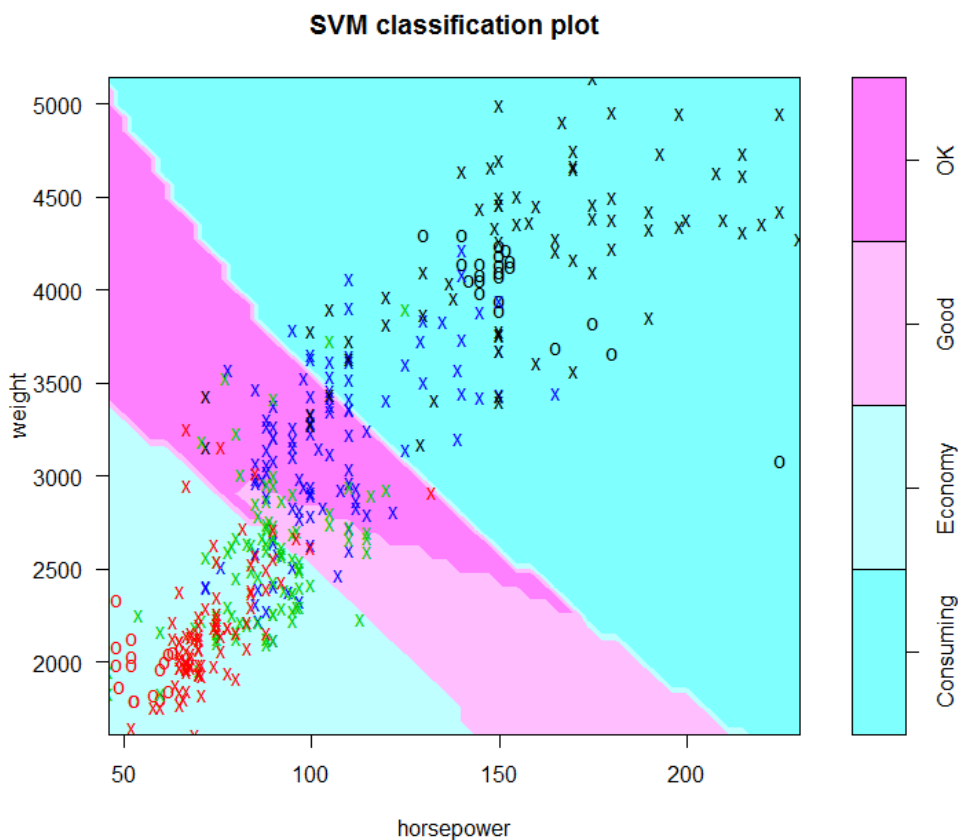
```
> ECO4[mpg < 22.75] = "OK"
> ECO4[mpg < 17] = "Consuming"
```

```
> table(ECO4)
ECO4
Consuming    Economy    Good    OK
      92         103      93    104
```

```
> S4 = svm( ECO4 ~ weight + horsepower, data=d, cost=0.1, kernel="sigmoid" )
Error in svm.default(x, y, scale = scale, ..., na.action = na.action) :
  Need numeric dependent variable for regression.
```

R was trying to do regression SVM but realized that ECO4 is not numerical. We can direct R to do classification by replacing ECO4 with factor(ECO4).

```
> S4 = svm( factor(ECO4) ~ weight + horsepower, data=d, cost=0.1, kernel="sigmoid" )
> plot(S4, data=d)
```



```
> Yhat = predict( S4, data.frame(Auto) )
> table( Yhat, ECO4 )
```

	ECO4			
Yhat	Consuming	Economy	Good	OK
Consuming	88	0	2	33
Economy	0	96	58	15
Good	0	2	9	5
OK	4	5	24	51

```
> mean( Yhat == ECO4 )
[1] 0.622449
```

It's more difficult to predict finer classes correctly